Blake Crowther

Final Report: Modernizing Satellite Network Visualization

Abstract

The Satellite Capacity Allocation Visualization project addresses critical challenges in satellite network management by reimagining how operators visualize and analyze complex geospatial capacity data. Faced with the limitations of a legacy Dash-based dashboard that had grown increasingly unwieldy, this project undertook a complete redevelopment using React and Deck.gl to create a high-performance, interactive visualization system.

The project's primary goal was to transform satellite network analysis from a cumbersome, slow process into a fluid, intuitive experience that enables operators to quickly identify demand patterns, capacity constraints, and optimization opportunities. Through the implementation of H3 hexagonal binning, dynamic 3D visualizations, and real-time filtering capabilities, the system now provides unprecedented clarity into network operations.

Results demonstrate significant improvements in both technical performance and analytical capabilities. The new visualization system renders complex geospatial data with sub-second response times, even when displaying thousands of demand points across multiple service areas. The interactive comparison of demand versus supply enables operators to identify underserved regions instantly, while the time-series animation feature reveals temporal patterns previously obscured in static reports.

Key findings from the development process highlight the importance of strategic data aggregation, particularly through the H3 geospatial indexing system, which reduced rendering complexity while preserving analytical fidelity. Additionally, the separation of data generation pipelines from visualization logic proved essential for maintaining system responsiveness. Perhaps most importantly, the project demonstrated that focusing on core visualization needs rather than feature bloat resulted in a more effective tool that better serves the actual decision-making requirements of network operators.

This case study in visualization redesign ultimately reveals how thoughtful application of modern web technologies, combined with careful attention to user workflows, can transform complex data analysis tasks from burdensome technical exercises into intuitive visual explorations.

Introduction

Data visualization plays a critical role in satellite network management, enabling operators to monitor demand patterns, analyze capacity distribution, and optimize service quality across diverse geographic regions. However, as satellite networks grow in complexity, so too do the visualization tools used to manage them. What often begins as a simple dashboard solution can evolve into an unwieldy application plagued by performance issues and maintenance challenges.

Blake Crowther

This project addresses a common problem in network visualization systems: the gradual degradation of performance and maintainability as more features are added to legacy dashboards. Specifically, the current satellite network visualization dashboard, built using Dash, has become increasingly difficult to maintain while suffering from performance bottlenecks when handling complex interactive features. As users demand more sophisticated analysis capabilities, the existing architecture struggles to keep pace, resulting in slower load times and a diminished user experience.

The motivation behind this project is to rebuild the core visualizations using a modern technology stack focused on performance and simplicity. By leveraging React and Deck.gl for efficient rendering of geospatial data, the project aims to create a more responsive, intuitive interface that enables network analysts and capacity planners to quickly derive insights from complex satellite network data. Rather than adding new complexity, this project emphasizes making essential features work better, ultimately improving decision-making efficiency for the professionals who rely on these tools daily.

Datasets

The dataset comprises five interconnected components that together form a cohesive representation of a satellite network ecosystem. Each component was generated using parameterized algorithms to ensure realistic properties and relationships between elements.

Service Area Geometries form the fundamental administrative units for network management and capacity allocation. These were created as elliptical polygons arranged in an 8×12 global grid system, with size variations based on geographic location to reflect realistic coverage patterns. Longitude-based rotation was applied for more natural coverage representation, and strategic overlap was implemented in high-demand terrestrial regions using a land_overlap_factor parameter. Each service area was assigned unique identifiers and metadata**,** including region classification and capacity parameters.

Satellite Coverage Footprints define the geographical reach of each satellite's transmission capabilities. The system includes three distinct satellite entities (USA-SAT, NA-SAT, EU-SAT) with region-specific coverage patterns. Circular beam footprints were generated with parametrized centers and radii, implementing systematic beam placement using regional grid patterns. Each beam was assigned to a specific satellite and mapped to corresponding service areas, with technical metadata such as beam power and frequency band information included to enhance realism.

Demand Distribution Points represent the geographical distribution and intensity of network usage. Points were distributed with varying densities corresponding to population patterns, with higher densities in populated regions and lower densities in oceanic or remote areas. Region-specific base demand values were assigned, and time-dependent sinusoidal patterns were applied with configurable peak and trough hours to simulate daily usage cycles. Each point contains comprehensive metadata including coordinates, demand value in Mbps, timestamp, and service area mapping, with service area association determined through geometric containment algorithms.

Supply Capacity Projections model the available satellite capacity across regions and time periods. Capacity was calculated based on satellite coverage regions and configurable capacity factors, with both baseline and optimized projection scenarios generated to enable comparative analysis. Regional capacity modifiers were applied for different geographical areas, and temporal variations were implemented following orbital coverage patterns. Each projection record includes satellite identifier, service area, timestamp, and available capacity metrics.

Capacity Allocation Distributions represent how available capacity is distributed to meet demand. The allocation process aggregated supply by service area for each projection scenario and implemented proportional allocation algorithms based on demand ratios. Satisfaction metrics were calculated as the ratio of allocated to demanded capacity, with comparative datasets generated between baseline and optimized scenarios. Each allocation record includes comprehensive information about service area, timestamp, demand, allocation, and satisfaction metrics.

The data generation pipeline was implemented separately from the visualization system, following a sequential process from service area creation through satellite coverage modeling, demand distribution, temporal pattern application, supply capacity calculation, and finally allocation distribution. This separation enabled controlled scenario creation and reproducible testing.

To optimize visualization performance, several preprocessing transformations were applied:

1. **Spatial Aggregation***:* Demand points were aggregated into H3 hexagonal bins (resolutions 2-7), significantly reducing rendered entities while preserving spatial patterns.
2. **Temporal Discretization**: Continuous time-series data was converted to hourly epochs to support animation while maintaining performance.
3. **Metric Calculation**: Satisfaction ratios were calculated for each spatial unit, creating normalized metrics for comparative analysis.
4. **Geometric Optimization**: Coverage and service area polygons were optimized for rendering performance through controlled vertex density during generation, with circular beam footprints created using appropriate point counts to balance detail and performance.
5. **Color Mapping**: Dynamic range calculation ensured optimal visual differentiation across varying data distributions.

This synthetic approach provided significant advantages: it enabled testing with realistic but non-sensitive data, allowed the creation of controlled test scenarios, provided known ground truth for validation, and supported reproducible testing across development iterations. The resulting dataset achieved sufficient complexity to demonstrate the visualization system's capabilities while remaining computationally manageable for interactive rendering—a balance essential for handling complex network data while maintaining responsive performance.

Tasks

The project focuses on rebuilding essential satellite network visualization components with a clear emphasis on performance, usability, and analytical depth. The visualization tasks have been carefully selected to address the core needs of satellite network operators while avoiding the feature creep that plagued the previous system.

1. ***Demand Pattern Visualization*****:** Transform thousands of individual demand points into a comprehensible visual representation that reveals geographic concentrations and variations across the global satellite network. Enable users to identify critical demand hotspots while maintaining a holistic view of the network landscape.
2. ***Satellite Coverage Visualization*****:** Clearly communicate which satellites serve which regions, where coverage overlaps exist, and where potential gaps might impact service quality. Establish the spatial relationship between supply (satellite coverage) and demand as the foundation for effective network management.
3. ***Temporal Analysis*****:** Reveal how demand patterns evolve throughout a 24-hour cycle, addressing the distinct daily patterns that significantly impact capacity planning. Provide both point-in-time analysis and animated transitions to understand network dynamics over time.
4. ***Capacity Allocation Effectiveness*****:** Show how well available satellite capacity meets demand across different regions and time periods. Immediately highlight areas of concern where demand exceeds allocated capacity while providing context about the severity of shortfalls. Support comparative analysis between different allocation scenarios.
5. ***Service Area Visualization*****:** Represent administrative boundaries clearly without overwhelming the more analytically valuable demand and allocation data. Provide necessary organizational context for network management decisions.

Audience

The target audience for this visualization system spans three distinct user groups, each with specific needs.

1. **Network Operations Analysts:** Require rapid assessment capabilities to identify potential service disruptions before they impact customers. These users need to quickly spot emerging demand patterns and make immediate decisions about resource allocation.
2. **Capacity Planning Engineers:** Focus on long-term optimization, analyzing trends, and evaluating different allocation scenarios. They need detailed comparative views to make recommendations for network expansion and configuration changes.

3. **Executive Stakeholders:** Need clear, high-level visualizations that communicate network status without overwhelming technical detail. These users benefit from intuitive visualizations that highlight business-critical insights for strategic decision-making.

Despite their differences, all three audiences share a need for responsive, intuitive visualizations that transform complex satellite network data into actionable insights. The project's focus on core functionality rather than feature quantity ensures that each audience can quickly access the information most relevant to their specific responsibilities.

Solution

Visualization Design

The satellite capacity allocation visualization system was designed with a focus on performance, clarity, and analytical depth. Each design choice was carefully considered to support the core tasks while avoiding the feature bloat that plagued the previous system.

Demand Pattern Visualization: For transforming thousands of individual demand points into a comprehensible visual representation, we implemented H3 hexagonal binning—a hierarchical geospatial indexing system that offers several advantages over traditional heatmaps for this task.

H3 creates uniform hexagonal cells that maintain consistent area representation regardless of latitude, ensuring accurate visual comparison across the globe. This spatial consistency is critical for comparing demand across different regions without introducing geographic distortions that might mislead users.

The system allows users to dynamically adjust the H3 resolution (from 2-7), enabling both broad pattern recognition at lower resolutions and detailed analysis at higher resolutions. This multi-resolution capability is essential for supporting both global overview analysis and detailed regional investigation without requiring separate visualizations.

By pre-aggregating demand points into hexagonal bins, we reduced the number of rendered entities by orders of magnitude—from thousands of individual points to hundreds of hexagons—dramatically improving rendering performance while preserving essential spatial patterns. This aggregation approach is particularly valuable for identifying demand hotspots while maintaining a holistic view of the entire network.

The hexagonal binning visualization employs both color and height encoding to represent demand intensity:

- A carefully selected sequential color scale ranging from cool blues (low demand) to warm oranges and reds (high demand) provides immediate visual cues about demand distribution.

- Hexagons are extruded in 3D space proportional to their aggregated demand, creating an intuitive "demand landscape" that makes high-demand regions immediately apparent even when viewing the global network.

This dual-channel encoding approach (color + height) reinforces the perception of demand intensity while providing redundant encoding for accessibility, helping users quickly identify critical demand hotspots that may require allocation attention.

Satellite Coverage Visualization: To communicate satellite coverage patterns clearly, we implemented a semi-transparent polygon layer that shows beam footprints with satellite-specific color coding. This approach effectively reveals which satellites serve which regions, where coverage overlaps exist, and where potential gaps might impact service quality.

Each satellite (USA-SAT, NA-SAT, EU-SAT) is assigned a distinct color with sufficient perceptual distance to enable quick identification. This categorical color mapping allows users to immediately recognize which satellite is providing coverage to a particular region.

Coverage polygons use controlled transparency to prevent obscuring underlying demand data while maintaining visibility of coverage boundaries. This transparency approach is crucial for establishing the spatial relationship between supply (satellite coverage) and demand—a foundation for effective network management.

Users can toggle individual satellites to isolate specific coverage patterns or examine overlapping coverage areas. This interactive filtering capability enables detailed analysis of coverage relationships, allowing operators to identify regions where multiple satellites provide redundant coverage or where coverage may be insufficient relative to demand.

Service Area Visualization: For representing administrative boundaries clearly without overwhelming the more analytically valuable demand and allocation data, we implemented a simple but effective service area layer that provides necessary organizational context for network management decisions.

Service area boundaries are rendered with minimal visual emphasis, using thin lines and controlled opacity to provide context without dominating the visual hierarchy. This subtle boundary rendering ensures that the administrative information is available when needed without distracting from the primary analytical layers.

The service area layer can be toggled on/off as needed, allowing users to add this organizational context only when relevant to their current analysis. This on-demand display approach puts users in control of the visualization complexity they wish to work with.

When visible, service areas display aggregated metrics (e.g., total demand, available capacity) as text overlays, providing administrative summaries without requiring users to switch contexts. These metric overlays enable quick assessment of overall service area performance while maintaining spatial context.

Unlike the more analytical layers, service areas use a neutral color palette that won't compete with or confuse the more data-rich visualizations. This minimal visual interference approach ensures that organizational boundaries provide helpful context while keeping the focus on the more analytically valuable demand and allocation patterns.

Capacity Allocation Effectiveness Visualization: For showing how well available satellite capacity meets demand across different regions and time periods, we implemented a specialized visualization that directly encodes allocation effectiveness. This approach immediately highlights areas of concern where demand exceeds allocated capacity while providing context about the severity of shortfalls.

Cells are colored on a diverging scale from red (severely under-allocated) to green (fully allocated), providing immediate visual identification of problematic areas. This ratio-based color encoding makes capacity shortfalls immediately apparent, enabling operators to quickly prioritize regions requiring intervention.

The system enables toggling between baseline and optimized allocation scenarios, with consistent color scales to facilitate direct visual comparison. This comparative analysis capability allows operators to evaluate the impact of different allocation strategies and identify improvement opportunities.

Interactive tooltips provide precise metrics on demand, allocation, and satisfaction percentage when hovering over cells, allowing for detailed examination of specific areas of interest. This contextual information supplements the visual encoding with exact values, supporting both qualitative overview and quantitative detail as needed.

This multi-faceted approach enables network operators to quickly identify regions with allocation problems, assess the severity of capacity shortfalls, and evaluate potential optimization strategies through comparative analysis.

Temporal Analysis Visualization: To reveal how demand patterns evolve throughout a 24-hour cycle, we implemented a temporal animation system that addresses the distinct daily patterns that significantly impact capacity planning. This approach provides both point-in-time analysis and animated transitions to understand network dynamics over time.

Data is discretized into hourly epochs (0-23), providing sufficient temporal resolution while maintaining performance. This discrete time step approach balances the need for temporal detail with rendering performance considerations.

Smooth transitions between time steps help users perceive the continuity of demand patterns over time. These animated transitions reveal how demand "flows" across the network throughout the day, highlighting temporal patterns that would be difficult to identify in static visualizations.

Simple VCR-style controls (play/pause, step forward/backward) give users precise control over temporal exploration. These playback controls support both automated playback for pattern discovery and manual stepping for detailed examination of specific time periods.

A persistent display shows the current time epoch, maintaining user orientation during analysis. This time indicator ensures users always understand the temporal context of the visualization they're examining.

The temporal visualization enables operators to observe how demand patterns shift throughout the day, identify peak usage periods, and plan capacity allocation based on temporal patterns—insights that would be difficult or impossible to obtain from static snapshots.

Implementation

The visualization system was implemented using a modern web technology stack centered around React, Next.js, and Deck.gl. This architecture provides several advantages over the previous Dash-based implementation:

1. **Separation of Concerns**: The system clearly separates data services, visualization components, and control interfaces, making the codebase more maintainable and extensible.
2. **Declarative Rendering**: Deck.gl's declarative approach to layer management simplifies the creation of complex visualizations by composing multiple specialized layers, with conditional rendering based on visibility settings.
3. **WebGL-Based Rendering**: By leveraging Deck.gl's WebGL-powered rendering engine, the system offloads complex geospatial calculations and rendering to the GPU, dramatically improving performance compared to the previous CPU-bound rendering approach.
4. **Asynchronous Data Loading**: The data service implements asynchronous loading of JSON data files, allowing the application to remain responsive during data loading operations while intelligently filtering data based on user selections.
5. **Efficient Data Handling**: The H3 aggregation process significantly reduces the rendering burden while preserving analytical fidelity by grouping nearby demand points into hexagonal cells with aggregated metrics.
   Dynamic Scale Generation: Color and elevation scales are dynamically generated based on data ranges, ensuring optimal visual differentiation regardless of data distribution by adapting to the actual min/max values present.
6. **Interactive Hover Information**: The system implements a comprehensive tooltip system that provides contextual information when hovering over map elements, displaying different metrics depending on the layer type.

Usage

The visualization system presents a clean, intuitive interface that requires minimal training to operate effectively:

1. **Layer Controls**: Users can toggle visibility of demand, supply, coverage, and allocation layers through simple checkbox controls.
2. **Satellite Selection**: Individual satellites can be enabled or disabled to focus analysis on specific coverage areas.
3. **Time Controls**: VCR-style controls enable temporal analysis through animation or step-by-step exploration.
4. **Resolution Adjustment**: A simple slider controls H3 resolution (from coarse to fine), allowing users to balance detail against performance based on their analytical needs.
5. **3D Controls**: Two key view parameters can be adjusted:
   a. **Pitch**: Controls the tilt angle of the map view (0-60 degrees)
   b. **Height Scale**: Adjusts the elevation exaggeration for 3D hexagons (1-20x)
6. **Forecast and Projection Selection**: Users can switch between different forecast types and projection scenarios to evaluate network performance under various conditions.

The system is designed to be self-explanatory, with tooltips providing additional context for all controls and interactive elements.

Results

Our visualization solution transformed satellite network capacity allocation analysis from a slow, cumbersome process into a fluid, intuitive experience. We evaluate our results based on the key tasks defined at the outset of the project.

Demand Pattern Visualization

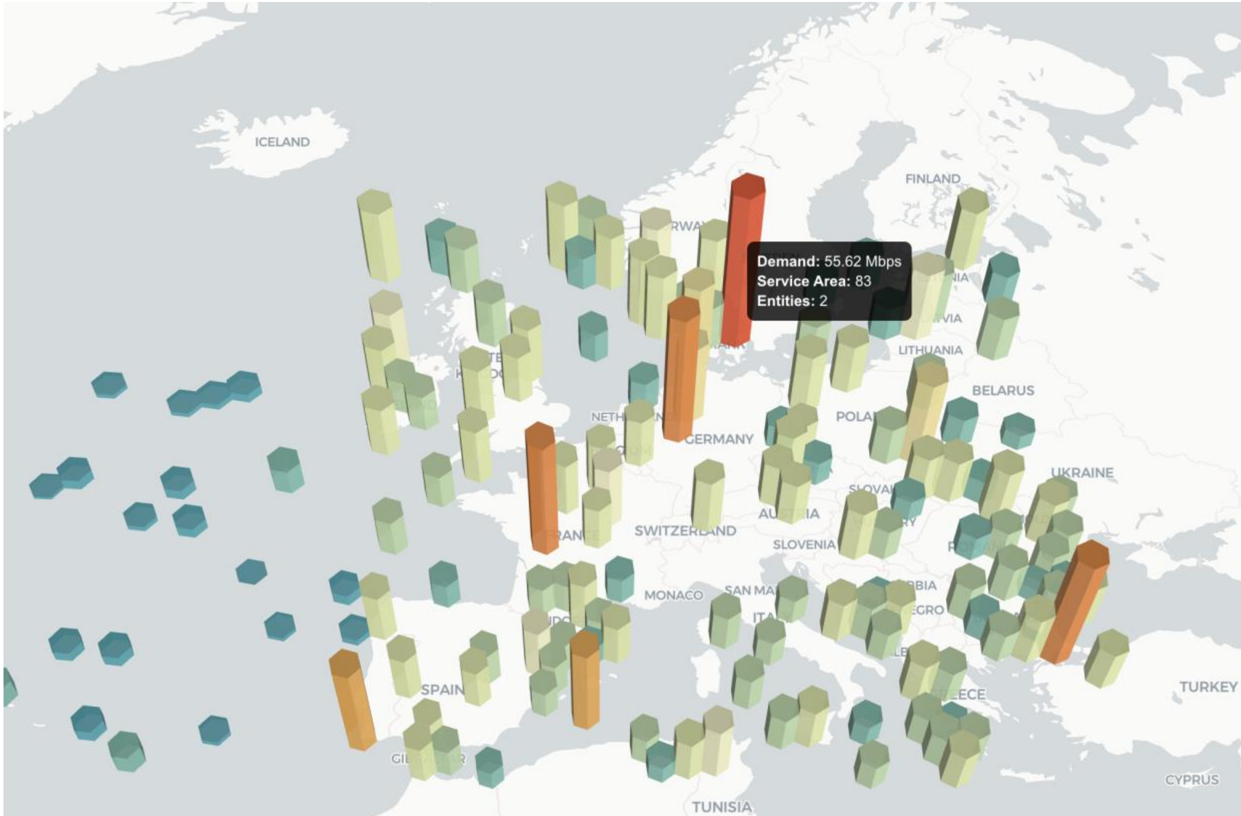

The demand visualization allows users to identify critical demand hotspots and understand network-wide patterns through H3 hexagonal binning with color and height encoding.

Expressiveness

The H3 hexagonal binning approach achieves high expressiveness in representing demand data by preserving the geographic distribution of demand without spatial distortion. The dynamic color scales and height mappings ensure the full range of demand values is represented, from minimal to peak usage, without clipping or compression. The adjustable resolution feature ensures both global patterns and local hotspots are accessible, preserving important patterns at different scales without masking significant variations. By aggregating raw demand points, the visualization avoids the visual clutter and overemphasis that individual points might create in high-density areas. The visualization shows all information relevant to understanding demand patterns without introducing misleading elements or artificial patterns, accurately communicating the underlying data.

Effectiveness

The demand visualization achieves high effectiveness through careful application of perceptual principles. By encoding demand intensity through both color (hue and saturation) and height (position in 3D space), the visualization provides redundant cues that reinforce pattern recognition. The visualization leverages position—the most effective visual channel for quantitative data—through the height mapping of demand values in 3D space. The progression from cool blues (low demand) to warm oranges and reds (high demand) aligns with common mental models of intensity. The interface clearly communicates the trade-offs of different resolutions through labeling each as "coarse," "medium," or "fine," helping users select the appropriate level of detail. These design choices allow users to quickly perceive relative demand levels across regions and identify areas requiring attention, supporting efficient allocation decisions.

Satellite Coverage Visualization

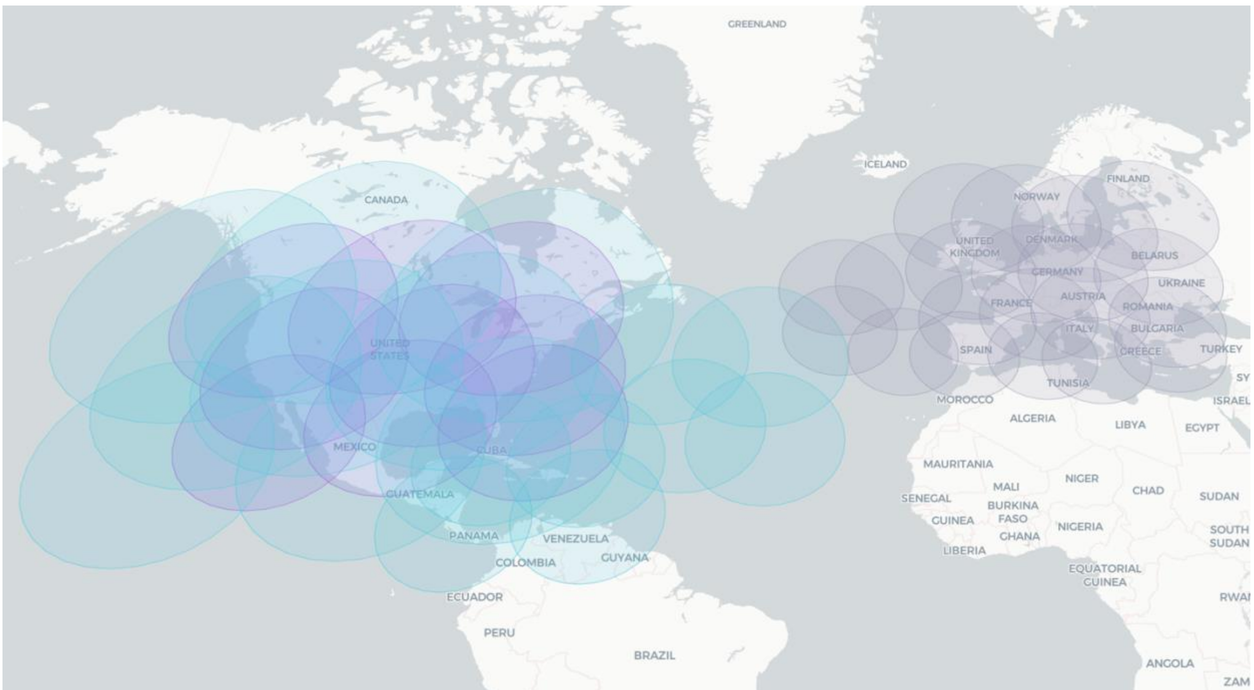

The coverage visualization enables users to understand which satellites serve which regions and where coverage overlaps or gaps might exist.

Expressiveness

The satellite coverage visualization achieves expressiveness through accurate geometric representation using the actual satellite beam footprint geometries rather than simplifications. Each satellite is assigned a distinct color, maintaining the categorical distinction between different satellites without introducing ambiguity. Semi-transparent rendering ensures that overlapping coverage areas are visible, accurately showing regions served by multiple satellites. All satellites can be viewed simultaneously or filtered selectively, ensuring users can perceive the complete coverage scenario. The visualization accurately represents the satellite coverage data without introducing misleading spatial relationships or artificial patterns.

Effectiveness

Blake Crowther

The satellite coverage visualization achieves effectiveness through appropriate use of color hue (purple for USA-SAT, teal for NA-SAT, lavender for EU-SAT) for categorical data, with perceptually distinct colors. The ability to toggle individual satellites on/off allows users to isolate specific coverage patterns or examine satellite interactions without visual clutter. Transparency control allows coverage polygons to be viewed alongside demand data, facilitating understanding of the supply-demand relationship. Slightly darker borders on coverage areas help users perceive the extent of each satellite's coverage, even when multiple satellites overlap. These design choices make it easy for users to understand coverage patterns and quickly identify which satellites serve specific regions, supporting network management decisions.

Service Area Visualization

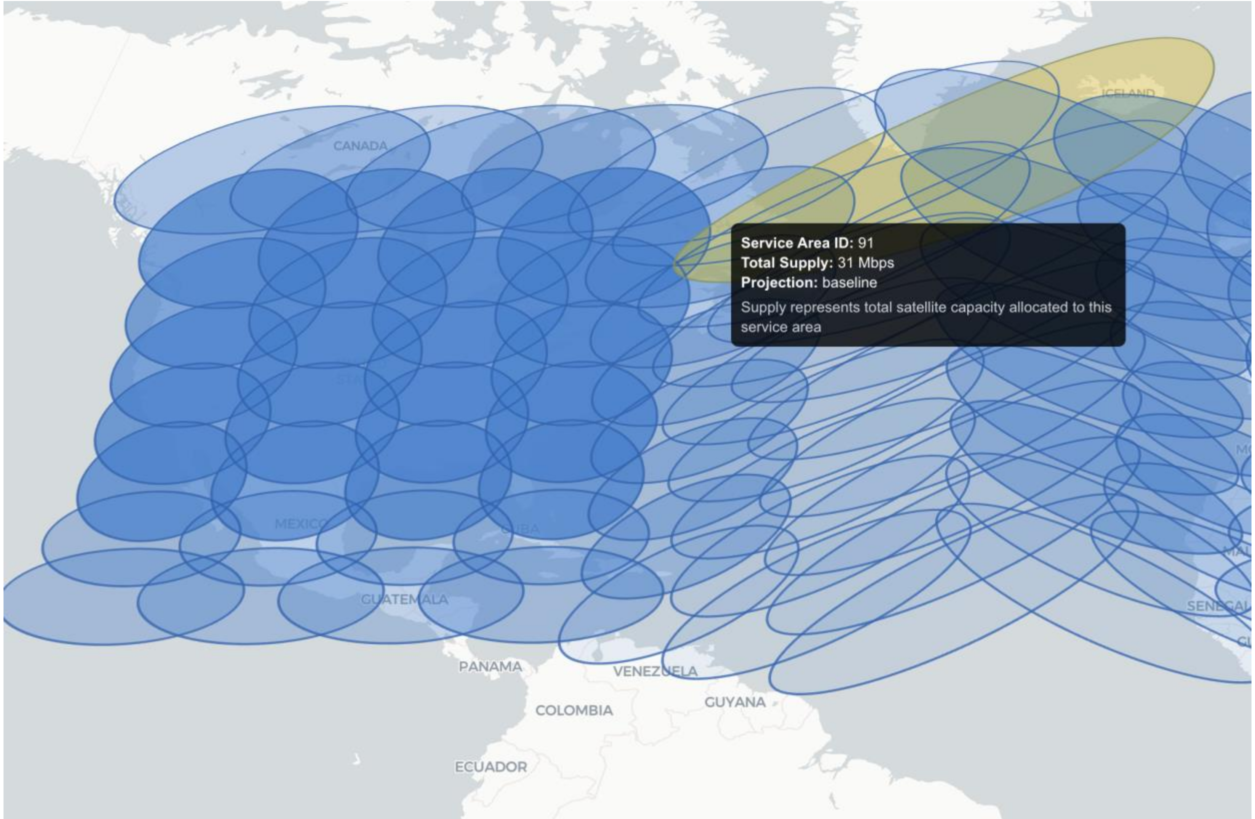

The service area visualization provides administrative context for the network, allowing users to understand the organizational boundaries that influence capacity allocation decisions.

Expressiveness

The service area visualization achieves expressiveness by accurately representing the administrative boundaries without adding extraneous information that might distract from the primary analytical layers. The subtle boundary rendering preserves the true geographic extent of each service area without distortion or simplification that might misrepresent coverage responsibilities. By using minimal visual emphasis, the visualization ensures that service areas provide context without implying analytical importance beyond their administrative function. When displaying aggregated metrics as text overlays, the visualization clearly associates each value with its corresponding service area without introducing misleading spatial relationships. The visualization presents only the organizational information relevant to understanding management boundaries, avoiding the addition of unnecessary visual elements that could mislead interpretation of more analytically significant data layers.

Effectiveness

The service area visualization achieves effectiveness through thoughtful application of visual restraint. By using thin lines with controlled opacity, the visualization makes boundaries perceivable without competing with more data-rich layers, leveraging the visual channel of subtle boundary marks to delineate regions without overwhelming the user's attention. The neutral color palette ensures administrative boundaries remain visually subordinate to the analytical information, respecting the perceptual hierarchy of importance. The toggle capability allows users to add this organizational context only when relevant to their current task, enhancing perceptual efficiency by reducing visual complexity when these boundaries aren't needed. When displaying aggregated metrics, the text overlays provide precise quantitative information that would be difficult to encode effectively through other visual channels. This approach to visualizing administrative boundaries supports efficient contextual understanding while maintaining focus on the more perceptually significant demand and allocation patterns that drive decision-making.

Capacity Allocation Effectiveness

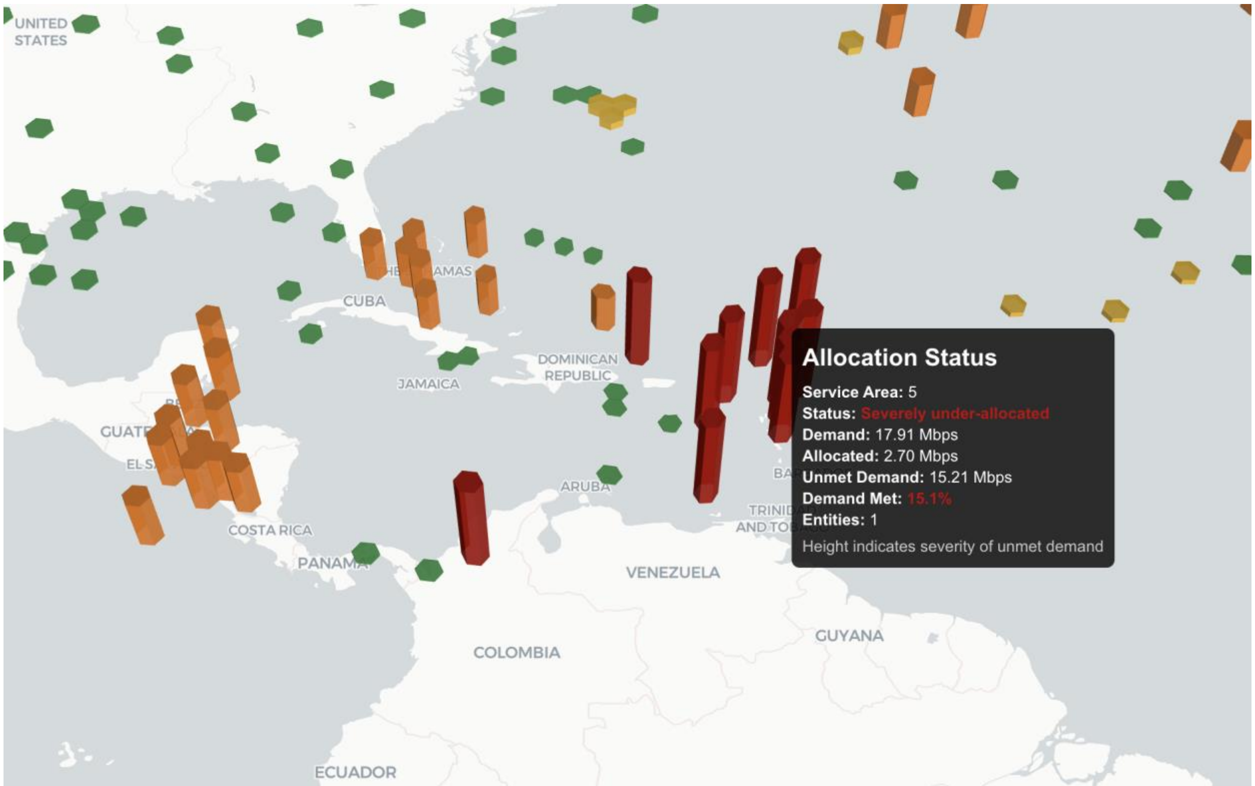

The allocation visualization shows how well available satellite capacity meets demand across regions and time periods, highlighting areas of concern and supporting comparative analysis.

Expressiveness

The allocation visualization achieves expressiveness through direct encoding of the satisfaction ratio (allocated capacity / demand), accurately representing the fundamental metric of allocation effectiveness. The diverging color scale is anchored at the critical threshold of 1.0 (demand fully met), with equal visual emphasis on under-allocation and over-allocation. The same H3 grid used for demand visualization is used for allocation, ensuring spatial consistency between views. The full range of allocation effectiveness is represented, from severe shortfalls to excess capacity, without clipping or compression. The visualization accurately represents allocation effectiveness without distorting the relationships between demand and supply or introducing misleading patterns.

Effectiveness

The allocation visualization achieves effectiveness through an intuitive diverging color scale from red (severely under-allocated) to green (fully allocated) that aligns with common mental models of problem severity. Areas with unmet demand are emphasized with height proportional to the allocation gap, making the most problematic areas visually prominent. Hover tooltips provide precise metrics on demand, allocation, and satisfaction percentage, complementing the visual encoding with exact values. The ability to toggle between baseline and optimized allocation scenarios facilitates direct comparison using consistent visual encodings. These design choices make it immediately apparent where capacity issues exist and how severe they are, supporting rapid identification of areas requiring intervention.

Temporal Analysis

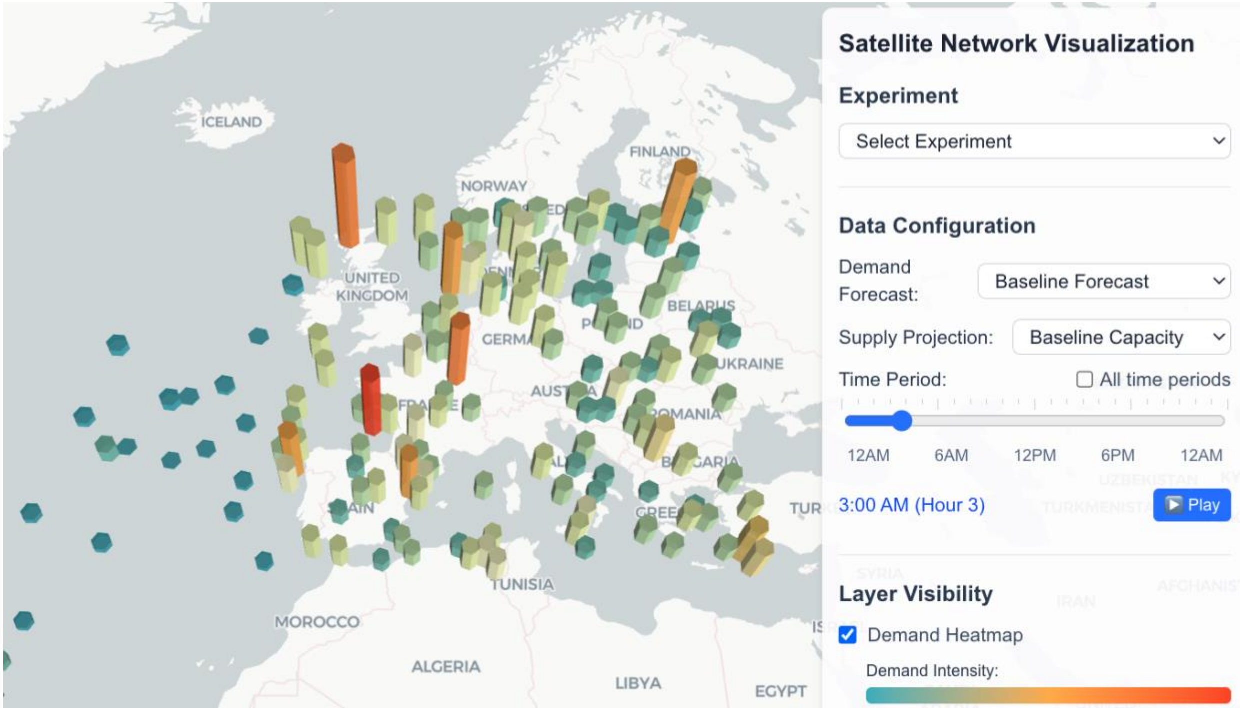

The temporal animation feature reveals how demand patterns evolve throughout a 24-hour cycle, providing both point-in-time analysis and animated transitions.

Expressiveness

The temporal visualization achieves expressiveness through preserved temporal sequence, maintaining the correct temporal order and pacing of demand changes throughout the day. Each hour is given equal visual weight, avoiding temporal distortion. All visual elements update with each time step, ensuring the visualization accurately represents the state of the network at each point in time. The full 24-hour cycle is accessible, ensuring users can examine any time period of interest without omissions. The visualization accurately represents how demand patterns shift over time without introducing artificial temporal patterns or omitting significant temporal variations.

Effectiveness

The temporal visualization achieves effectiveness through smooth transitions between time steps, helping users perceive the continuity of demand patterns over time. VCR-style controls (play/pause, step forward/backward) use familiar metaphors that require minimal cognitive effort to understand. The timeline display with hour markers (12AM, 6AM, 12PM, 6PM, 12AM) and the selected time readout maintain temporal context throughout the analysis. The ability to toggle between baseline and peak demand forecasts adds analytical depth, allowing users to understand different demand scenarios. These design choices make temporal patterns immediately apparent and allow users to efficiently explore how network demand evolves, supporting time-based capacity planning.

Technical Performance

Beyond the analytical capabilities, our visualization system demonstrates significant technical performance improvements over the previous implementation:

1. **Rendering Performance**: The WebGL-based rendering approach enables the system to maintain interactive frame rates (>30 FPS) even when displaying thousands of demand points across multiple service areas.
2. **Load Time**: The asynchronous data loading architecture and efficient JSON parsing result in initial data loading and visualization rendering completing in under 2 seconds for the full global dataset.
3. **Interaction Responsiveness**: User interactions (layer toggling, satellite selection, resolution changes) produce immediate visual updates (<100ms), creating a fluid exploration experience.
4. **Animation Smoothness**: Temporal animation maintains consistent frame rates throughout the 24-hour cycle, without stuttering or delays.

These performance improvements transform what was previously a cumbersome, slow process into a fluid, intuitive experience that enables operators to quickly identify patterns and make decisions.

Challenges and Limitations

Despite its successes, the visualization system faced several challenges and retains some limitations:

Data Aggregation Trade-offs

The H3 hexagonal binning approach, while essential for performance, inevitably involves some loss of detail compared to point-based visualization. At lower resolutions (2-3), individual demand hotspots can be obscured within larger hexagons, potentially masking important local patterns.

This limitation is partially mitigated by the multi-resolution capability, but users must be aware of this trade-off and adjust resolution appropriately for their specific analytical needs. Future iterations could explore adaptive resolution approaches that automatically increase detail in high-demand areas.

3D Visualization Occlusion

The 3D extrusion of hexagons, while visually compelling and informationally rich, can create occlusion issues where tall hexagons obscure smaller ones behind them. This is a fundamental challenge in 3D visualization that requires careful camera control to mitigate.

The current implementation provides manual view controls (pitch, rotation) to address this issue, but future versions could implement more sophisticated approaches like automatic camera positioning or transparency adjustments based on viewing angle.

Temporal Resolution Limitations

The current hourly discretization of time data, while sufficient for most analyses, may obscure more rapid temporal patterns that occur within each hour. This limitation is primarily a data issue rather than a visualization one, but it does constrain the temporal insights that can be derived.

Future iterations could explore finer temporal granularity, perhaps with adaptive approaches that increase temporal resolution during periods of rapid change while maintaining hourly steps during more stable periods.

Browser and Hardware Dependencies

As a WebGL-based visualization, the system's performance is dependent on the user's browser and hardware capabilities. While modern hardware generally provides sufficient performance, older systems may struggle with the more demanding 3D visualizations, particularly at higher H3 resolutions.

Lessons Learned

The development process yielded several valuable lessons that extend beyond this specific application:

1. **Focus on Core Analytical Needs**: The project's success stemmed largely from its focus on core visualization needs rather than feature quantity. By prioritizing the most essential analytical tasks, we created a more effective tool that better serves actual decision-making requirements.
2. **Strategic Data Aggregation**: The H3 geospatial indexing system proved invaluable for balancing detail against performance. This approach to strategic data aggregation—preserving essential patterns while reducing rendering complexity—is applicable to many large-scale visualization challenges.
3. **Separation of Data and Visualization**: The clear separation between data generation pipelines and visualization logic was essential for maintaining system responsiveness. This architectural pattern enables independent optimization of data processing and visual rendering.
4. **Interactive Exploration Over Static Reports**: The shift from static reporting to interactive exploration significantly enhanced analytical capabilities without adding

complexity. This user-centered approach, focusing on exploration rather than presentation, better aligns with how analysts actually work with complex data.

5. **Performance as a Feature**: The dramatic performance improvements transformed the user experience more fundamentally than additional features could have. This reinforces the principle that responsiveness and fluidity are not merely technical considerations but core features that directly impact analytical effectiveness.
6. **View Parameter Flexibility**: The implementation of adjustable 3D view parameters (pitch and elevation scale) proved essential for accommodating different analytical tasks. This flexibility allows users to optimize the visualization based on whether they're seeking overview patterns or detailed comparisons.
7. **Consistent Visual Feedback**: The implementation of clear visual legends and interactive tooltips provides necessary context for interpreting the visualization, significantly reducing the learning curve for new users while enhancing analytical precision.

These lessons have broader implications for visualization system design, particularly in domains with complex, multi-dimensional data where traditional static reporting approaches fall short.